

Benchmark Design Concepts

Contents

- Basic concepts of benchmarking
- Processor benchmarks
- Disk benchmarks

Conditions that benchmark programs must satisfy

- Reliable work in all environments
- Same workload in all environments
- Reproducibility of results
- No machine dependent components

Reliable work in various environments

- Use algorithms that are sufficiently tested, reliable and generate stable results
- Benchmark programs must satisfy highest portability standards
- Benchmark programs should perform same operations in all hardware and software environments

Stability of workload

- Benchmark workload must be the same for all competitive systems:
 - All systems must use the same algorithms
 - All systems must use the same data
- Benchmark programs should not use machine-dependent or OS-dependent components

Reproducibility of results

- Benchmark programs are regularly executed multiple times to compute the average run time, or to adjust desired run time
- In each execution benchmark program must generate same results, i.e. results must be reproducible

Compilation of benchmark programs

- Avoid using different compilers for different computer systems or different operating systems.
- Make sure to always compile benchmark programs using the same (highest) level of optimization
- Avoid using debug versions of programs for benchmarking

Selection of benchmark workloads

- In the case of specific natural workload that benchmarks must simulate, benchmark workload is selected to be as similar to the natural workload as possible
- If information about natural workload is not available, we use **default workload** that includes operations that are frequent in all cases of data processing

Default workloads

- Default processor workload contains two groups of operations:
 - Floating-point operations
 - Integer and combinatorial operations
- Default disk workload includes two groups of operations
 - Sequential write/read of a disk file
 - Random write/read of a disk file

Speed indicators

- If n benchmark programs B_1, \dots, B_n run for T_1, \dots, T_n seconds, then the computer speed in operations per minute can be defined as $V_1 = 60/T_1, \dots, V_n = 60/T_n$
- The overall (average) speed indicator V is computed as the harmonic mean of individual speed indicators:

$$\begin{aligned} V &= n / (1/V_1 + \dots + 1/V_n) \\ &= 60n / (T_1 + \dots + T_n) \end{aligned}$$

Processor benchmarks

The size of benchmark

- If a benchmark program is large and/or processes large data sets, then it will use processor, cache memories, bus, and the main memory
- Small benchmark programs that can fit in cache memory may execute from a cache memory with negligible use of bus and main memory

Typical floating point benchmarks

- Matrix inversion and other matrix operations
- Solution of linear algebraic equations
- Roots of polynomials
- Numeric integration
- Computation of functions

Matrix inversion benchmark

- Rule #1: Avoid using matrices that contain random numbers – such matrices are frequently singular or can cause other numerical problems
- Use matrices that are proved to be regular.
- Strictly diagonally dominant matrices are non-singular

Strictly diagonally dominant matrix

- Strictly diagonally dominant matrix satisfies the condition

$$|a_{ii}| > \sum_{i \neq j} |a_{ij}|, \quad i = 1, \dots, n, \quad j = 1, \dots, n$$

- This matrix is nonsingular (invertible), and therefore it is convenient for benchmarking
- There are other matrices that are also known as nonsingular

Invertible matrices that can cause numeric problems in benchmarking

- Cauchy's matrix

$$a[i][j] = 1 / (i + j)$$

- Hilbert's matrix

$$a[i][j] = 1 / (i + j - 1)$$

- Vandermonde's matrix

$$a[i][j] = j^{i-1}$$

An invertible combinatorial matrix that is suitable for benchmarking

$$\begin{aligned}a[i][j] &= 1.0001, \quad i \neq j \\ &= 2.0001, \quad i = j \quad (\text{any size})\end{aligned}$$

2.0001	1.0001	1.0001	1.0001	1.0001	1.0001	1.0001	1.0001
1.0001	2.0001	1.0001	1.0001	1.0001	1.0001	1.0001	1.0001
1.0001	1.0001	2.0001	1.0001	1.0001	1.0001	1.0001	1.0001
1.0001	1.0001	1.0001	2.0001	1.0001	1.0001	1.0001	1.0001
1.0001	1.0001	1.0001	1.0001	2.0001	1.0001	1.0001	1.0001
1.0001	1.0001	1.0001	1.0001	1.0001	2.0001	1.0001	1.0001
1.0001	1.0001	1.0001	1.0001	1.0001	1.0001	2.0001	1.0001
1.0001	1.0001	1.0001	1.0001	1.0001	1.0001	1.0001	2.0001

Sample
matrix of
size $n=8$

Single index matrices

- In benchmarking it is frequently necessary to adjust the size of matrix to achieve a desired run time or other properties
- Matrices with 2 indices are not suitable for frequent changes of size
- Single-index matrices are just a single array and their size is easily adjusted. They are convenient for benchmarking.

Avoid using library functions

- Library functions from `math.h`, such as `sin`, `cos`, `exp`, `log`, etc. differ from compiler to compiler, and therefore may be a source of differences in seemingly equal workloads
- Another function that is also machine-dependent is `rand()` from `stdlib.h`
- Library functions should be used in benchmarking only when natural workloads use library functions with high frequency

Typical integer/combinatorial benchmarks

- Sorting arrays
- Search algorithms
- Random number generation and testing
- Combinatorial optimization
- Puzzles (e.g. 8 queens)
- Recursive algorithms

What data are convenient for sorting in benchmark programs?

- Random numbers can be convenient for sorting provided that they are fully reproducible on all machines starting with those having 16 bit integers.
- Deterministic arrays can also be suitable provided that they contain some regular nontrivial pattern

Random number generators for benchmarking

- Random number generators for benchmarking must satisfy the following conditions:
 - Same behavior on machines having the word size of 16 bits and more
 - Long cycle of random numbers
 - Simplicity and speed
 - Quality of distribution of random numbers is not important

Fibonacci-style generators

- $N(i) = (N(i-1) + N(i-2)) \bmod M, \quad i=2,3,\dots$
- Very simple and fast
- We must select $N(0)$, $N(1)$, and M
- Cycle repeats when
$$N(c)=N(1) \text{ and } N(c-1)=N(0)$$

The cycle size is c .
- If working with 16-bit integers the condition for avoiding overflow is
$$N(i-1) + N(i-2) < 2^{14} = 16384$$

What are the best values for $N(0)$, $N(1)$, and M ?

- Theoretical answer to this question is not simple [Knuth, Vol. 2]
- Empirical answer to this question is rather simple and based on experiments with a program that computes the maximum cycle length for various values of $N(0)$ and $N(1)$
- In the following program we use $p=N(i-2)$, $q=N(i-1)$, $r=N(i)$


```

#include<iostream.h>
#include<stdlib.h>
#include<iomanip.h>

void main(void)  // Finding maximum cycle of a Fibonacci-style random number generator
{
    unsigned long int cycle=0, cycle_max=0;
    unsigned int M_init=10000, M, N0=1, N1=2, M_opt, p, q, r;
    for(int k=0; k<20; k++)
    {
        for(M=M_init; M<16384; M++)  // 2^14=16384 or 2^15=32768
        {
            p=N0; q=N1; cycle=0;
            do
            {
                r = (p+q) % M; cycle++;
                p = q;  q = r;
            }while (!(p == N0) && (q == N1));
            if(cycle>cycle_max) {cycle_max = cycle; M_opt = M;}
        }
        cout << "N0=" << setw(5) << N0 << "    N1=" << setw(5) << N1
              << "    M_opt=" << setw(6) << M_opt
              << "    Max_cycle=" << setw(6) << cycle_max << "\n";
        N0=rand()%M_init; N1=rand()%M_init;
    }
}

```

N0= 1	N1= 2	M_opt= 15625	Max_cycle= 62500
N0= 41	N1= 8467	M_opt= 15625	Max_cycle= 62500
N0= 6334	N1= 6500	M_opt= 15625	Max_cycle= 62500
N0= 9169	N1= 5724	M_opt= 15625	Max_cycle= 62500
N0= 1478	N1= 9358	M_opt= 15625	Max_cycle= 62500
N0= 6962	N1= 4464	M_opt= 15625	Max_cycle= 62500
N0= 5705	N1= 8145	M_opt= 15625	Max_cycle= 62500
N0= 3281	N1= 6827	M_opt= 15625	Max_cycle= 62500
N0= 9961	N1= 491	M_opt= 15625	Max_cycle= 62500
N0= 2995	N1= 1942	M_opt= 15625	Max_cycle= 62500
N0= 4827	N1= 5436	M_opt= 15625	Max_cycle= 62500
N0= 2391	N1= 4604	M_opt= 15625	Max_cycle= 62500
N0= 3902	N1= 153	M_opt= 15625	Max_cycle= 62500
N0= 292	N1= 2382	M_opt= 15625	Max_cycle= 62500
N0= 7421	N1= 8716	M_opt= 15625	Max_cycle= 62500
N0= 9718	N1= 9895	M_opt= 15625	Max_cycle= 62500
N0= 5447	N1= 1726	M_opt= 15625	Max_cycle= 62500
N0= 4771	N1= 1538	M_opt= 15625	Max_cycle= 62500
N0= 1869	N1= 9912	M_opt= 15625	Max_cycle= 62500
N0= 5667	N1= 6299	M_opt= 15625	Max_cycle= 62500

Optimum
modulus and
maximum
cycle size for
20 random
seeds in the
case of
integers
where
 $M < 16384$

N0=	1	N1=	2	M_opt=	31250	Max_cycle=	187500
N0=	41	N1=	8467	M_opt=	31250	Max_cycle=	187500
N0=	6334	N1=	6500	M_opt=	31250	Max_cycle=	187500
N0=	9169	N1=	5724	M_opt=	31250	Max_cycle=	187500
N0=	1478	N1=	9358	M_opt=	31250	Max_cycle=	187500
N0=	6962	N1=	4464	M_opt=	31250	Max_cycle=	187500
N0=	5705	N1=	8145	M_opt=	31250	Max_cycle=	187500
N0=	3281	N1=	6827	M_opt=	31250	Max_cycle=	187500
N0=	9961	N1=	491	M_opt=	31250	Max_cycle=	187500
N0=	2995	N1=	1942	M_opt=	31250	Max_cycle=	187500
N0=	4827	N1=	5436	M_opt=	31250	Max_cycle=	187500
N0=	2391	N1=	4604	M_opt=	31250	Max_cycle=	187500
N0=	3902	N1=	153	M_opt=	31250	Max_cycle=	187500
N0=	292	N1=	2382	M_opt=	31250	Max_cycle=	187500
N0=	7421	N1=	8716	M_opt=	31250	Max_cycle=	187500
N0=	9718	N1=	9895	M_opt=	31250	Max_cycle=	187500
N0=	5447	N1=	1726	M_opt=	31250	Max_cycle=	187500
N0=	4771	N1=	1538	M_opt=	31250	Max_cycle=	187500
N0=	1869	N1=	9912	M_opt=	31250	Max_cycle=	187500
N0=	5667	N1=	6299	M_opt=	31250	Max_cycle=	187500

Optimum
modulus and
maximum
cycle size for
20 random
seeds in the
case of
unsigned
integers where
 $M < 32768$

Machine-independent random number generator working with integers

```
int rng_i( )           // Random integers
{
    // Cycle size = 62500

    static int p=1, q=2, r;

    r = (p+q) % 15625; p=q; q=r;

    return r;
}
```

Machine-independent random number generator working with unsigned integers

```
unsigned int rng_ui( ) // Unsigned random integers
{
    // Cycle size = 187500

    static unsigned int p=1, q=2, r;

    r = (p+q) % 31250; p=q; q=r;

    return r;
}
```

Properties of rng_i() and rng_ui()

- Portable generator
- Reproducible Fibonacci-style random sequence
- Sequence length = 62500 or 187500
- No overflow even with 16-bit machines
- Suitable for sort benchmarks

Testing rng_ui and rng_i

- Generate 62500 random numbers with rng_i
- Generate 187500 random numbers with rng_i
- In both cases the last two numbers must be 1 and 2

```

int rng_i( )           // Random integers: cycle size = 62500
{
    static int p=1, q=2, r;
    r = (p+q)%15625; p=q; q=r;
    return r;
}

unsigned int rng_ui( ) // Unsigned random integers: cycle size = 187500
{
    static unsigned int p=1, q=2, r;
    r = (p+q)%31250; p=q; q=r;
    return r;
}

void main(void) // Testing rng_i and rng_ui
{
    for(int i=0; i<62498; i++) rng_i();
    cout << rng_i(); cout << ' ' << rng_i() << endl;

    for(i=0; i<187498; i++) rng_ui();
    cout << rng_ui(); cout << ' ' << rng_ui() << endl;
}

1 2
1 2

```


Deterministic sequences

- Deterministic sequences can be generated by repeating the integer expression $n = n+1 - n/M*M$. The resulting sequence is:
 $1, 2, \dots, M, 1, 2, \dots, M, \dots$
- Sample generator loop:
`for(k=0; k<kmax; k++) a[k]=n=n+1-n/M*M;`
- Such sequences can be sorted in decreasing order

Some simple algorithms that can be used as small CPU benchmarks

Linear Algorithm: Factorial

```
int f(int n)
{
    return (n<1) ? 1 : (n*f(n-1));
}
```

```
int F(int n)
{
    int f=1, i;
    for(i=2; i<=n; i++) f *= i;
    return f ;
}
```

Useful for very short run times

$T(n) = O(n)$ (both iterative and recursive)

Linear Algorithm: Two Variables

```
void merge (int a[ ], int na, int b[ ], int nb, int c[ ], int& nc)
{
    int i,j ;
    nc=i=j=0 ;
    while (i<na && j<nb) c[nc++] = (a[i]<b[j]) ? a[i++] : b[j++];
    while (i<na)         c[nc++] = a[i++] ;
    while (j<nb)         c[nc++] = b[j++] ;
}
```

Useful for very short run times

Run time: $T(na, nb) = c \cdot nc = c(na+nb)$

Logarithmic Algorithm: Binary Search

```
int bsearch(int v[], int n, int x)
{ int low, high, mid ;
  low=0 ; high = n-1 ;
  while (low <= high)
  { mid = (low + high) / 2 ;
    if      (x < v[mid])  high = mid-1 ;
    else if (x > v[mid])  low  = mid+1 ;
    else return mid ;
  }
  return -1 ; /* no match */
}
```

Useful for very short run times: $O(\log n)$. One of the most popular algorithms in Computer Science

Recursive Binary Search

```
int bsearch(int v[ ], int low, int high, int x)
{
    int mid = (low + high) / 2;

    if (low > high) return -1;

    if (x < v[mid]) return bsearch(v, low, mid-1, x);
    if (x > v[mid]) return bsearch(v, mid+1, high, x);

    return mid;
}
```

Useful for testing recursion.
Very short run times: $O(\log n)$.

Quicksort: $O(n \log_2(n))$

```
void sort(int a[], int left, int right)
{ int i, mid;

  if(left >= right) return;

  for(mid=left, i=left+1; i <= right; i++)
    if(a[i] < a[left]) swap(a, ++mid, i);

  Swap(v, left, mid);
  sort(v, left, mid-1);
  sort(a, mid+1, right);
}
```

One of the most popular algorithms in Computer Science. Frequently used and useful for benchmarking

Quadratic Algorithms: $O(n^2)$

For each component of data set process all components of the data set:

```
for (i=0; i<n; i++)  
    for (i=0; i<n; i++)  
        {Processing}
```

These algorithms may be useful for easy increase of run time by increasing n .

$$T(n) = c \cdot n \cdot n = O(n^2); \quad c = \text{const.}$$

Simple Select Sort

```
void sort( int a[], int n )  
{ int i, j;  
  for( i=0 ; i<n-1 ; i++ )  
    for( j=i+1 ; j<n ; j++ )  
      if( a[i] > a[j] ) swap(a[i], a[j]);  
}
```

Inefficient but extremely
simple algorithm.
Sometimes used for
simplistic benchmarking.

$$\begin{aligned} T(n) &\approx c(\text{number of executed if statements}) = c(1 + 2 + \dots + n-1) \\ &= cn(n-1)/2 = c(n^2-n)/2 \end{aligned}$$

$$T(n) = O(n^2)$$

Quadratic algorithm: Matrix Initialization

```
for (i=0; i<n; i++)  
    for (j=0; j<n; j++)  
        a[i][j] = rand( );
```

$$T(n) = c n^2 = O(n^2)$$

Quadratic algorithm: Matrix Transposition

```
for (i=0; i<n-1; i++)  
    for (j=i+1; j<n; j++)  
        swap(a[i][j] , a[j][i]);
```

$$T(n) = c(1+2+\dots+n-1) = c (n^2 - n) / 2 = O(n^2)$$

Matrix Multiplication

```
for (i=0; i<n; i++)
```

```
    for (j=0; j<n; j++)
```

```
        for (c[i][j] = k = 0; k<n; k++)
```

```
            c[i][j] += a[i][k]*b[k][j];
```

The most frequent matrix operation. Easy to increase the run time by adjusting n.

$$T(n) = c \cdot n \cdot n \cdot n = O(n^3); \quad c = \text{const.}$$

Exponential Algorithm: Fibonacci Numbers: 0, 1, 1, 2, 3, 5, 8, 13, ...

```
int f(int n)
{
    return (n<2) ? n : f(n-1)+f(n-2) ;
}
```

Extremely inefficient and almost illegal in programming practice. Can be used for testing recursion only.

$$\begin{aligned} T(n) &= t + T(n-1) + T(n-2) \approx T(n-1) + T(n-2), & n > 1 \\ &= 1, & n = 1 \\ &= 0, & n = 0 \end{aligned}$$

Solving $T(n) = T(n-1) + T(n-2)$

Suppose that $T(n) = cg^n$, $g > 0$

$$cg^n = cg^{n-1} + cg^{n-2}$$

$$1 = g^{-1} + g^{-2}$$

$$g^2 - g - 1 = 0$$

$$g = \frac{1 \pm \sqrt{1+4}}{2}; \quad g = \frac{1 + \sqrt{5}}{2} = 1.618$$

$$T(n) = c 1.618^n = O(1.618^n)$$

Linear Fibonacci Numbers

```
int f(int n)
{
    int i, zero=0, first=1, second=n;
    for(i=2; i<=n; i++)
    {
        second = first + zero;
        zero = first;  first = second;
    }
    return second;
}
```

This is a proper way to compute Fibonacci numbers. Insufficient complexity and too short run time make this program unsuitable for benchmarking.

$$T(n) = c \cdot n = O(n); \quad c = \text{const.}$$

How to make a SpeedMark benchmark


```
// Pseudocode of Speedmark benchmark (HW#1)
// 1. Run 10 seconds matrix inversion (or similar workload)
// 2. Compute matrix inversion (floating point) speed
// 3. Run 10 seconds quicksort (or similar workload)
// 4. Compute quicksort (integer) speed
// 5. Compute and display the average speed in operations/min
```

```
int NINT=0, NFLOAT=0;
```

```
double START, VINT, VFLOAT, AverageSpeed;
```

```
START = sec( );
```

```
while(sec( ) < START+10) {Minv( ); NFLOAT++;}
```

```
VFLOAT = 60*NFLOAT/(sec( )-START);
```

```
START = sec( );
```

```
while(sec( ) < START+10) {Qsort( ); NINT++;}
```

```
VINT = 60*NINT/(sec( )-START);
```

```
AverageSpeed = 2*VFLOAT*VINT/(VFLOAT+VINT);
```

```
Display VFLOAT, VINT, AverageSpeed;
```

```
// Select Minv( ) and Qsort( ) so that they run between 1 and 3 sec
```

Average speed as a harmonic mean

$VFLOAT$ = speed of floating point operations

$VINT$ = speed of integer operations

$$\begin{aligned} AVERAGE_SPEED &= \frac{2}{\frac{1}{VFLOAT} + \frac{1}{VINT}} \\ &= \frac{2 \times VFLOAT \times VINT}{VFLOAT + VINT} \end{aligned}$$

Disk benchmarks

Default workloads

- Writing a sequential file
- Reading a sequential file
- Random writing in a relative file
- Random reading from a relative file

Problems

- What is the available space on disk?
- How big should be the file?
- Contiguous or fragmented files?
- How to compare disks that have different capacities?

Disk speed as a harmonic mean

$VSEQ$ = speed of sequential read operations

$VRAN$ = speed of random read operations

$$\begin{aligned} AVERAGE_SPEED &= \frac{2}{\frac{1}{VSEQ} + \frac{1}{VRAN}} \\ &= \frac{2 \times VSEQ \times VRAN}{VSEQ + VRAN} \end{aligned}$$

Other disk operations

- Disk sequential write (SEQW)
- Disk sequential read (SEQR)
- Disk random write (RANW)
- Disk random read (RANR)

Disk speed as a harmonic mean

(a) The case of equal weights

$VSEQW$ = speed of sequential write operations

$VRANW$ = speed of random write operations

$VSEQR$ = speed of sequential read operations

$VRANR$ = speed of random read operations

$$AVERAGE_SPEED = \frac{4}{\frac{1}{VSEQR} + \frac{1}{VRANR} + \frac{1}{VSEQW} + \frac{1}{VRANW}}$$

Disk speed as a harmonic mean

(b) The case of different weights

$VSEQW$ = speed of sequential write operations

$VRANW$ = speed of random write operations

$VSEQR$ = speed of sequential read operations

$VRANR$ = speed of random read operations

$$AVERAGE_SPEED = \frac{1}{\frac{WSEQR}{VSEQR} + \frac{WRANR}{VRANR} + \frac{WSEQW}{VSEQW} + \frac{WRANW}{VRANW}}$$

$$WSEQR > 0, \quad WRANR > 0, \quad WSEQW > 0, \quad WRANW > 0$$


$$WSEQR + WRANR + WSEQW + WRANW = 1$$

Using different weights

- Weights determine relative importance
- Generally, the frequency of sequential and random disk operations is not the same
- Generally, the frequency of read and write disk operations is not the same
- Different weights can be used to take into account the different frequency of use


Review of basic disk operations in C++

Input File Stream




Input file stream is a sequence of bytes (characters) generated when we read data from a disk file. In C++ the input file stream can have any name: e.g., **IS**.

Output File Stream



Output file stream is a sequence of bytes (characters) generated when we write data to a disk file. In C++ the output file stream can have any name: e.g., **OS**.

FILES



Stream = sequence of bytes and a pointer showing the current byte (where insert/extract can start)

FILE = stream of characters terminated by the End Of File (EOF) record; also called *file stream*

Location: Disk, Floppy disk, tape, CD ROM, (sometimes even in memory), etc.

Size: Limited only by medium capacity

C++ FILE OPERATIONS

Library: `#include <fstream.h> // Support for
file stream operations`

Declaration: `ifstream IS; // Input file stream
ofstream OS; // Output file stream`

File name: (1) internal, and (2) external

Internal name: IS, and OS are **internal names**
used inside a C++ program to refer
to file streams. Same rules as for
any variable name.

External name: depends on rules for file names
used by the operating system

OPEN FILE: includes several operations:

- connect internal name IS and external name "data.txt"
- check access right of the user
- create buffer areas in memory
- transfer directory record from directory file to operating system

CLOSE FILE: includes several operations:


- disconnect internal name IS and external name "data.txt" (IS becomes an undefined variable)
- release buffer areas in memory
- transfer updated directory record from the OS to the directory file

General rule: OPEN as late as possible, and CLOSE as soon as possible, but do not exaggerate! Power failure or OS crash when a file is open can cause some data to be lost.

Purpose: File can be open for three purposes defined using input/output stream (ios) class member functions:

- read **ios::in**
- write **ios::out**
- append **ios::app**

READ:




IS.open(EFname, ios::in); // open for read

Declaration with initialization:

ifstream IS = EFname ; // shortest form, *for Borland compilers*
but NOT STANDARD
→ **ifstream IS(EFname, ios::in) ;**

WRITE:




```
OS.open(EFname, ios::out); // open for  
                           write
```

If the file **EFname** exists prior to execution of the **open()** function, then its contents is lost because EOF will be moved at the beginning of the file (**ios::out** is destructive!)

Declaration with initialization:

```
ofstream OS = EFname ; (NOT STANDARD !)  
→ ofstream OS(EFname, ios::out) ;
```

APPEND:



```
OS.open(EFname, ios::app); // open for  
                           append
```

EFname exists prior to execution of the **open()** function, and its contents is preserved

Declaration with initialization:

```
ofstream OS(EFname, ios::app) ;
```

EXTRACT (READ) DATA

```
int data;  
char c, s[80];
```

```
IS >> data ;
```

```
IS.getline(s, 80);  
IS.get(c);
```

- (1) Skip white spaces (blanks, tabs, newlines, formfeeds, etc.)**
- (2) Extract (and convert) data as long as possible**
- (3) Stop conversion when a white space or non-convertible character is encountered (this also includes the EOF record)**

CLOSE FILE:

```
IS.close() ; // IS becomes undefined
```


```
OS.close(); // OS becomes undefined
```

DETECT EOF:

```
IS >> data ;  
while(! IS.eof( ))  
{  
    // Process data  
    IS >> data ;  
}
```

Function **eof()** returns true (1) if EOF is detected

INSERT (WRITE) DATA



OS << data ;

OS << s ; // write a line into file

OS.put(c) ; // write one character

TEST OPEN FILE SUCCESS/FAILURE

Reason for failure:

- no file
- wrong directory
- no medium
- hardware error

TEST FUNCTION:

IS.fail()

OS.fail()

0 = successful last operation with IS/OS

!0 (usually 1) = last operation with IS/OS
failed

PASSING OF FILE STREAM ARGUMENTS (as other compound objects - only by reference)

```
void io(ifstream& IS, ofstream& OS, double& data)
{
    int a, b, c;
    IS >> a >> b >> c; // READ FROM IS
    OS << a << b;      // COPY TO OS
    data = a+b+c;      // RETURN VALUE
}                      // IS and OS are formal streams
                      (they do not exist as data objects)

void main(void)
{
    int result ;
    ifstream INPUT = "in.dat" ; // actual input stream
    ofstream OUTPUT = "out.dat" ; // actual out stream
    io(INPUT, OUTPUT, result) ; // Function call


    .....
    INPUT.close( ) ; // if close is omitted, it will
    OUTPUT.close( ) ; // be done automatically at
    .....          // the end of main program
}
```


Relative Files

- Relative files consist of indexed records
- Relative files support direct access (access according to the byte offset: the position is defined as the distance in bytes from the beginning of file)
- Each record can be designed as a struct

File model: a sequence of bytes

**Read/write pointer
positioned by seek oper.**



Record structure and size

- Example: struct client
 { int account;
 char name[20];
 double balance;
 } ;
- Size of record is measured in bytes and determined using the sizeof() function:
 sizeof(client)

Initialization of record

```
client c ;  
client blank = { 0, "", 0.}
```

```
cin >> c.account  
    >> c.name  
    >> c.balance ;
```

Basic operations: write

- Opening a relative file for writing:
ofstream OS(filename, ios::ate)
(ate = positioning “at end”)
- Writing records in arbitrary order:
 - Selecting file record (positioning file pointer)
OS.seekp(<byte offset>)
 - OS.write(<record pointer>, <record size>)
- seekp : positioning for put operation

Basic operations: read

- Opening a relative file for reading:
ofstream OS(filename, ios::in) (seq)
ofstream OS(filename, ios::ate) (random)
- Reading records in arbitrary order:
 - Selecting file record (positioning file pointer)
OS.seekg(<byte offset>)
 - OS.read(<record pointer>, <record size>)
- seekg : positioning for get operation

Cast operator (type)

- Conversion to desired data type
- Syntax:
 (desired data type) expression
- Cast is a unary operator and has higher precedence than arithmetic operations
- Example:
 $X = (\text{double}) (n + 1)$
- Alternative way:
 $X = \text{double} (n + 1)$

More casting

- `int k=1, n=2;`
`double x = (double) k/n // x = 0.5`
`// The first operation is the higher`
`// precedence unary (double)k`
`// Then n is automatically promoted to`
`// double`
- `(char *) &<struct record> // address of`
`an arbitrary record converted to a`
`character pointer`

Record pointer

- The type of record pointer is `char *`
- The address `<struct record>` has the type different from `char *`
- Conversion to type `char *` is performed using `cast (char *) <struct record>`
- `IS.read((char *) <rec>, sizeof(rec))`
- `OS.write((char *) <rec>, sizeof(rec))`